



Lightweighting DL Models

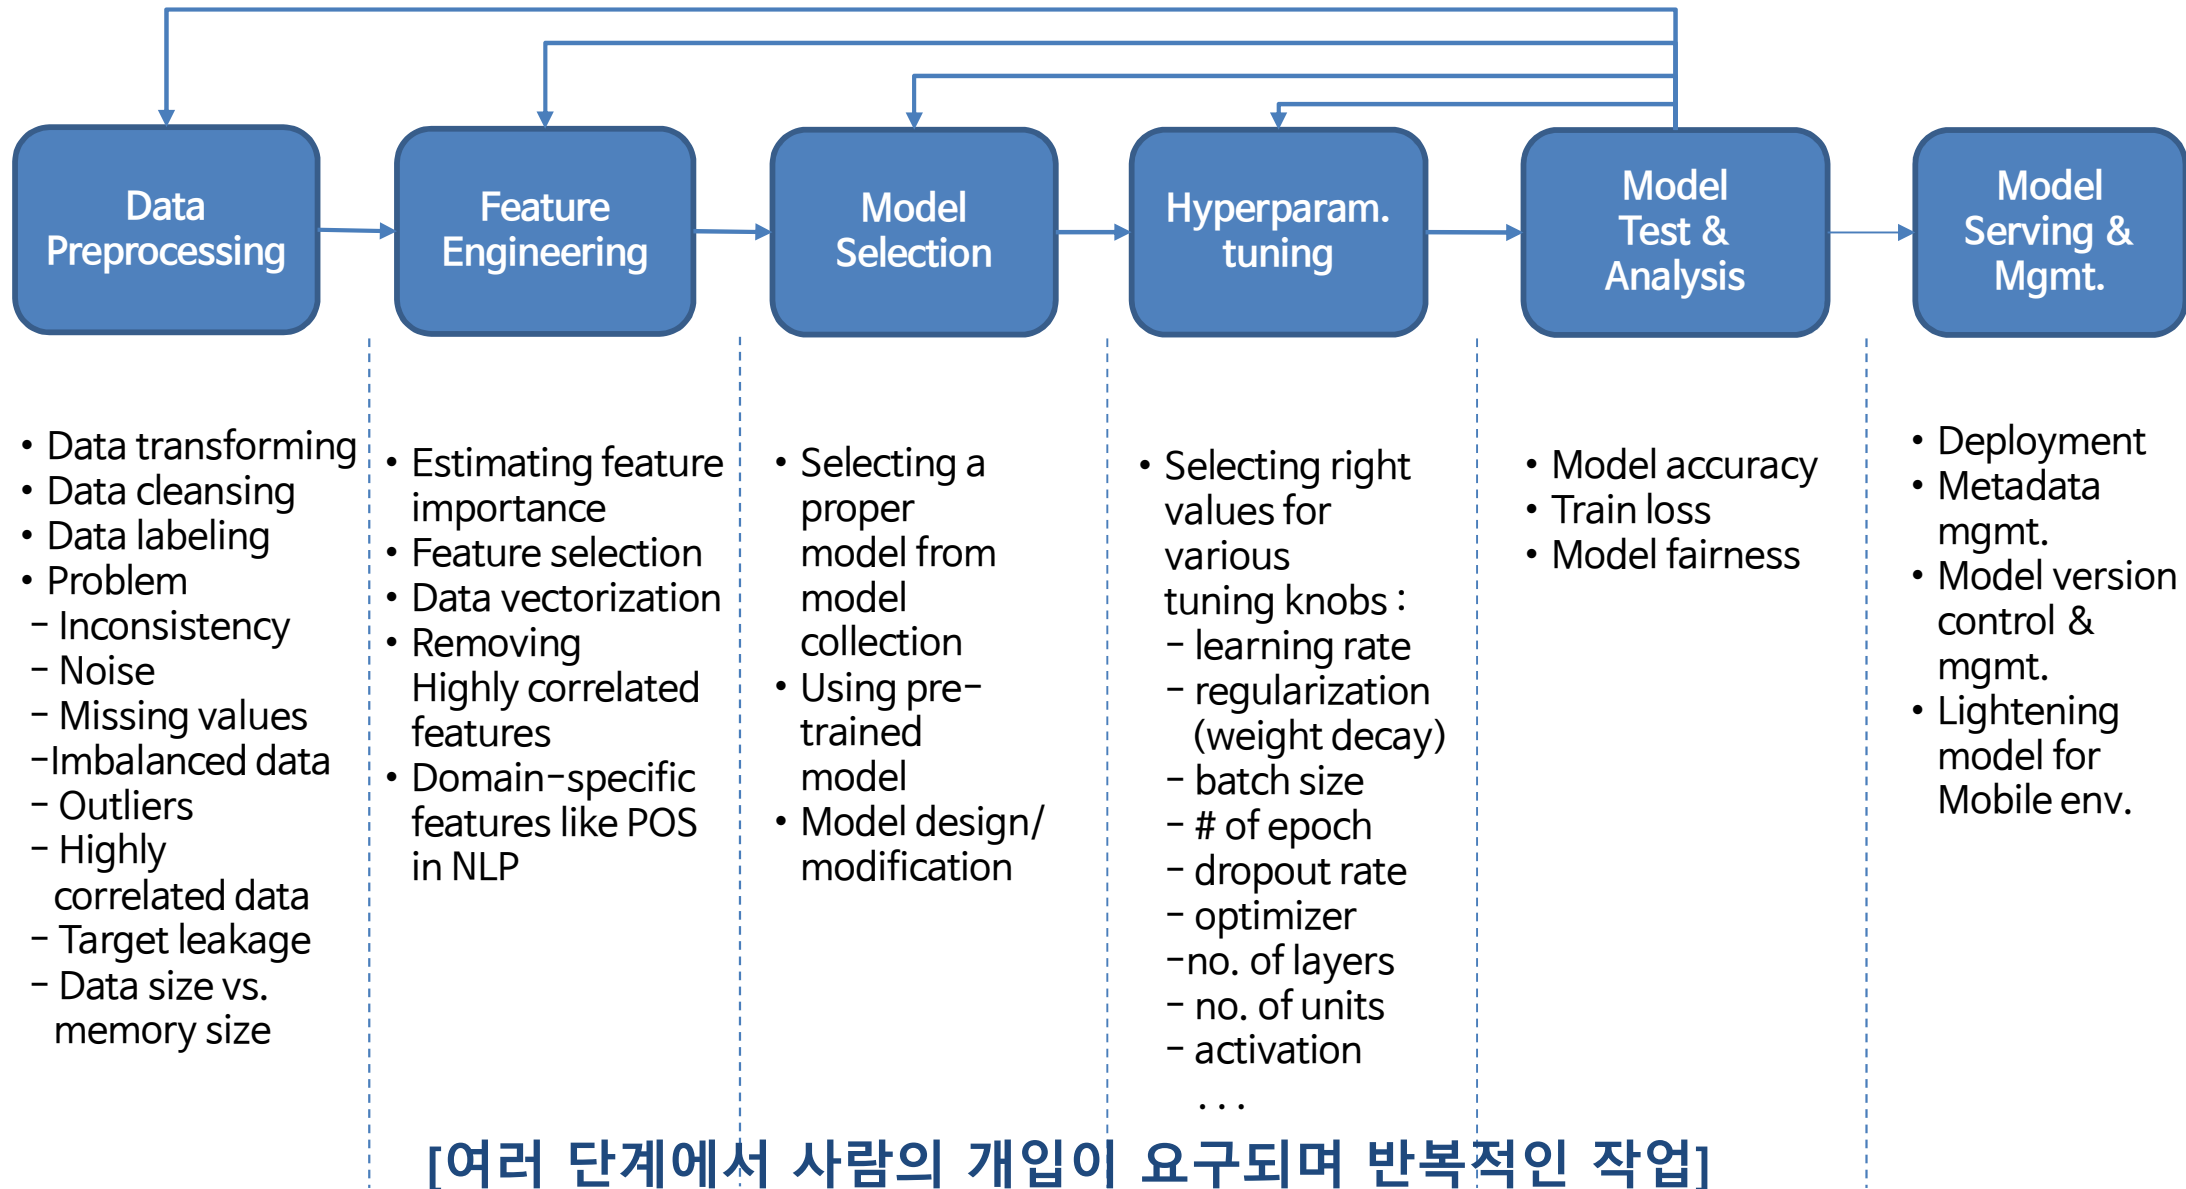
June 04-11, 2021
Kyong-Ha Lee
bart7449@gmail.com

I. DL process

- 대량의 Labeled data의 필요
 - 주로 지도 학습 기반으로 모델 훈련을 위해 대량의 정답 집합이 요구됨
- 반복적인 단계별 시행착오 과정의 연속
 - 개발자가 요구하는 정확도에 도달할 때까지 여러 단계에서 사람이 개입하면서 trial-or-error 방식으로 반복적인 훈련, 테스트 과정을 수행
 - **반복적인 Human-In-The-Loop 프로세스**
- 학습과 추론에 높은 수준의 전산 자원 요구
 - 많은 계층, 많은 파라미터 개수, 그에 따른 연산 수 및 메모리 증가
 - 예) BERT Large - 24 layers, 340M parameters, koBERT의 경우 V100 GPU x 32, horovod(w/ InfiniBand)로 학습에 약 한달 소요
- 학습 과정은 Batch 과정으로 수행
 - 데이터 변화에 따른 모델 재훈련 필요


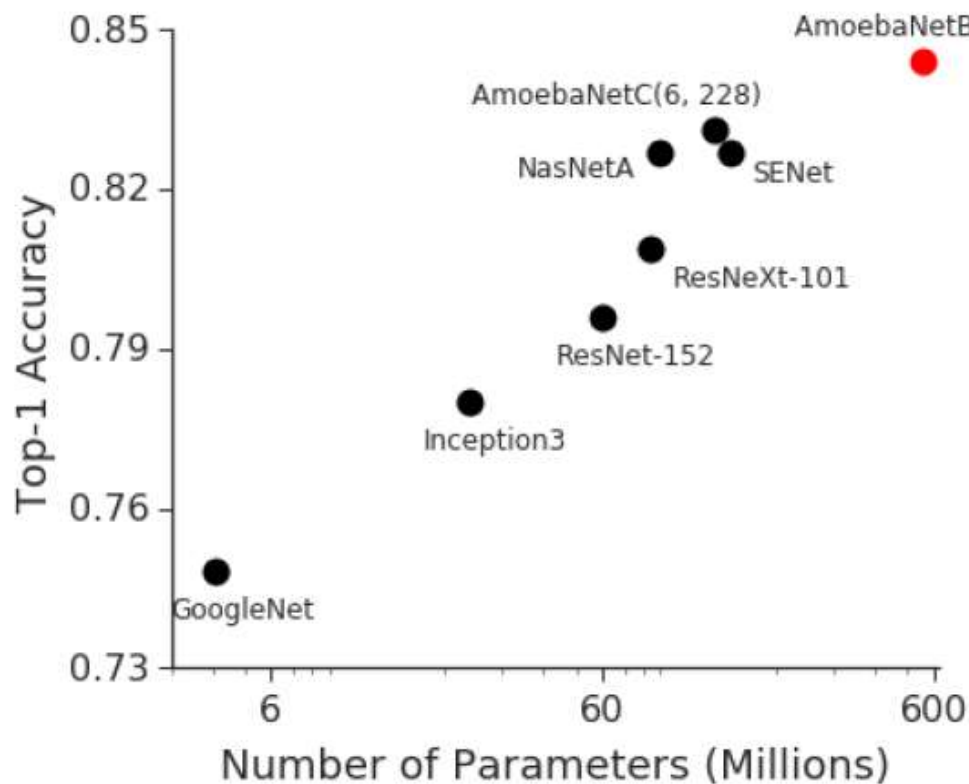


Challenge 1 : Human-In-The-Loop Process



Challenge 2: Model Size Matters!

- CNN-based image classifiers



Recent advances have shown that ever-larger DNN models lead to better task performance and past progress in visual recognition tasks has also shown **a strong correlation between the model size and accuracy**

* An analysis of Deep Neural Network Models for Practical Applications, A. Canziani et al., April 2017

** Benchmark analysis of representative deep neural network architecture, Blanco et al., Oct. 2018

Challenge 3: Time & Costs

Long training time limits ML researcher's productivity

Correlation btw. #layers and time

Model	Error rate	Training time
ResNet18	10.76%	2.5 days
ResNet50	7.02%	5 days
ResNet101	6.21%	1 week
ResNet150	6.16%	1.5 weeks

* M40 GPU, fb.resnet.torch

- KoBERT (SKT, Oct. 2019)
 - 24 layers, 340M parameters
 - **1 month** with **32 V100** GPUs interconnected with Horovod(w/ infiniband)
- XLNet (Yang, arXiv 19 Jun 2019)
 - 340 million parameters
 - Training : **2.5 days with 512 TPU v3 chips** for 500k steps
 - 512 TPU x 2.5 days x \$8 = **\$ 245,000**
- Gpipe (Huang, NIPS Dec. 2019)
 - 556 million parameters
- NASNet (Barret, CVPR June 2018)
 - 800 GPU, **28 days** training
- GPT-3 (OpenAI, 2020)
 - 175B parameters, required 3.14E23 FLOPS for training
 - At theoretical 28 TFLOPS for V100, **355 GPU-years** and cost \$4.6M for a single training run
 - 700GB memory to store it in FP32

Challenge 4: Energy Efficiency

- AlphaGo : 1,920 CPUs and 280 GPUs, \$3,000 electric bill per game
- “Training a single DL model can emit as much carbon as 5 cars in their lifetimes” – MIT Tech. Review, 2019

Common carbon footprint benchmarks

in lbs of CO2 equivalent

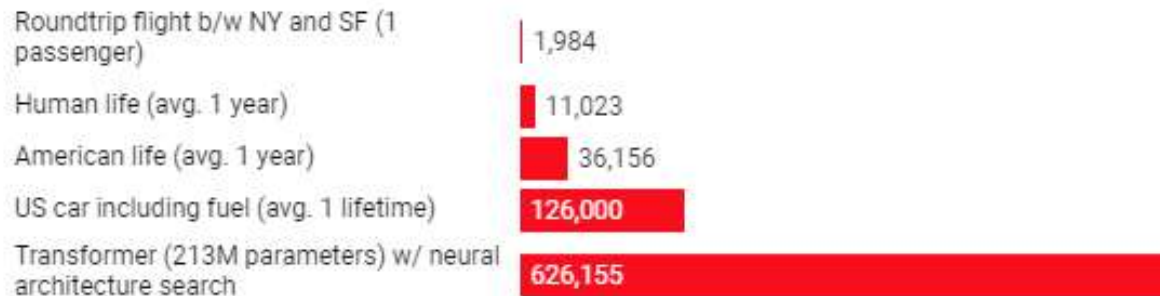


Chart: MIT Technology Review • Source: Strubell et al. • Created with Datawrapper

The estimated costs of training a model

	Date of original paper	Energy consumption (kWh)	Carbon footprint (lbs of CO2e)	Cloud compute cost (USD)
Transformer (65M parameters)	Jun, 2017	27	26	\$41-\$140
Transformer (213M parameters)	Jun, 2017	201	192	\$289-\$981
ELMo	Feb, 2018	275	262	\$433-\$1,472
BERT (110M parameters)	Oct, 2018	1,507	1,438	\$3,751-\$12,571
Transformer (213M parameters) w/ neural architecture search	Jan, 2019	656,347	626,155	\$942,973-\$3,201,722
GPT-2	Feb, 2019	-	-	\$12,902-\$43,008

Note: Because of a lack of power draw data on GPT-2's training hardware, the researchers weren't able to calculate its carbon footprint.

Table: MIT Technology Review • Source: Strubell et al. • Created with Datawrapper

- **DL development is naturally an Iterative HITL process**
 - Learning and test processes are performed in a batch job
 - Usually developed in a trial-and-error fashion
 - **Massive labeled data are required** for learning accurate DL models
- **Recent models become much larger**
 - Even a single model is composed of 175 billion parameters (i.e., GPT-3, ~652GBs)
 - All the **SIZE**, **TIME**, and **COST** grow very fast beyond our resources
 - Learning models tend to become harder with a limited budget and time
- **Expensive computational resources are required**
 - Energy efficiency is also an issue

II. 모델 경량화: Overview

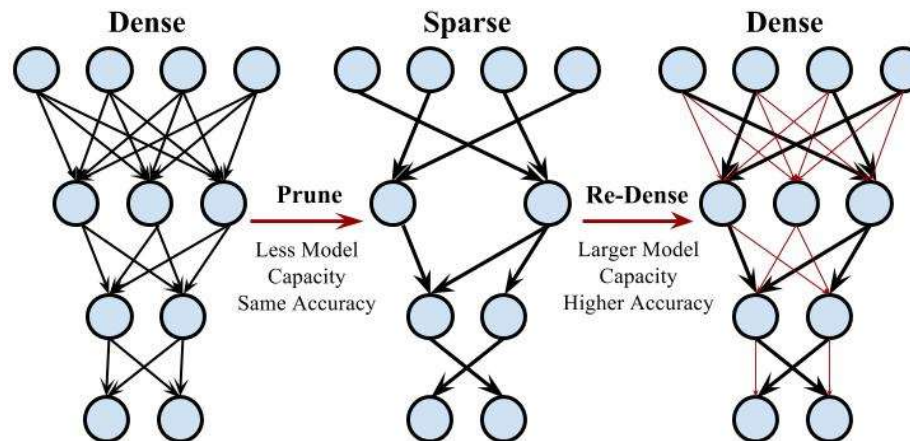
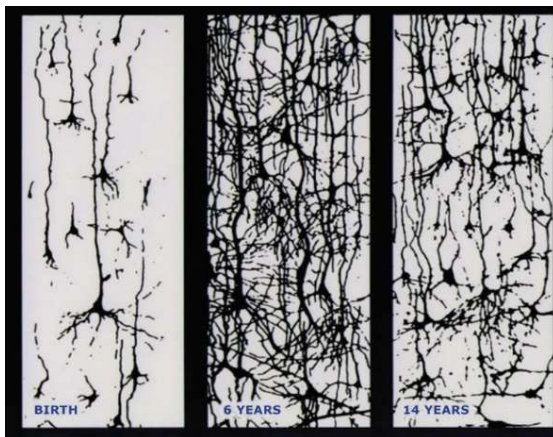
- 기술의 목적
 - 기존 full-strength 모델로부터 정확도 손실을 최소화하면서 보다 고속의 추론이 가능하고 메모리, 에너지 면에서 **효율적인 경량화된 모델**의 생성
 - 모바일 단말 등 저성능의 전산자원에서의 구동을 목적으로 함
 - 모델 압축(model compression)이라고도 함
- 모델 경량화 기법의 분류

Approaches	Technique	Approach
Pruning	• Remove weights or neurons under threshold	• From pre-trained model
Weight sharing	• Cluster weights and encode centroid of the weights	• From pre-trained model
Quantization	• Substitute weight, activations or gradients with lower bit-widths	• From pre-trained model • From scratch
Distillation	• Teacher-student model	• From pre-trained model
Low-rank approximation	• Tensor decomposition on convolutional tensor	• From pre-trained model
Sparse regularization	• Learn a structure-regularized version of CNN to achieve speed-up with little accuracy loss	• From pre-trained model
Compact Network Design	• Revise networks to be more computationally efficient	• From scratch

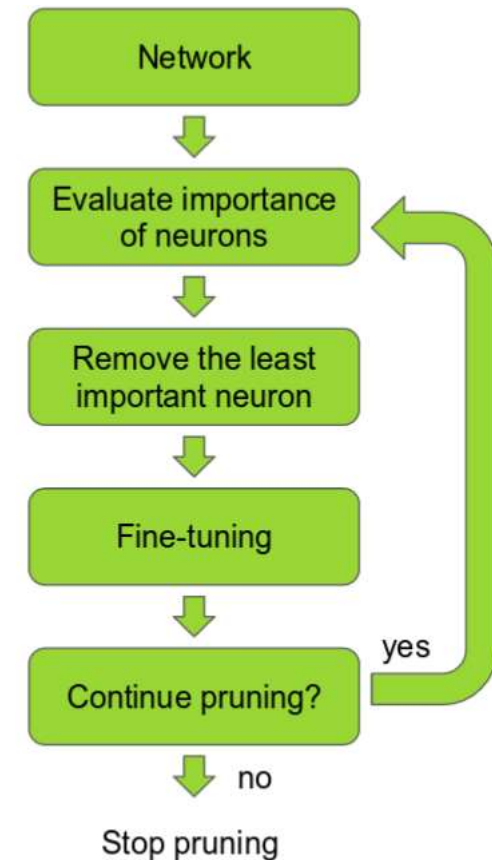
[모바일 단말에의 이식 및 추론을 위한 초기 목적]

II. 모델 경량화: Pruning

- Pruning unimportant neurons in DNN
 - Motivated by how real brain learns
 - Remove** weights which $|weight| < threshold$
 - Retrain** after pruning weights
 - Learn effective connections by iterative pruning



[기존 모델의 가공 및 재학습을 통한 모델 경량화]


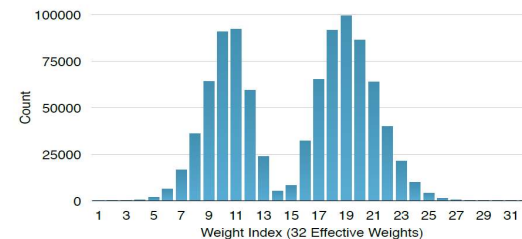
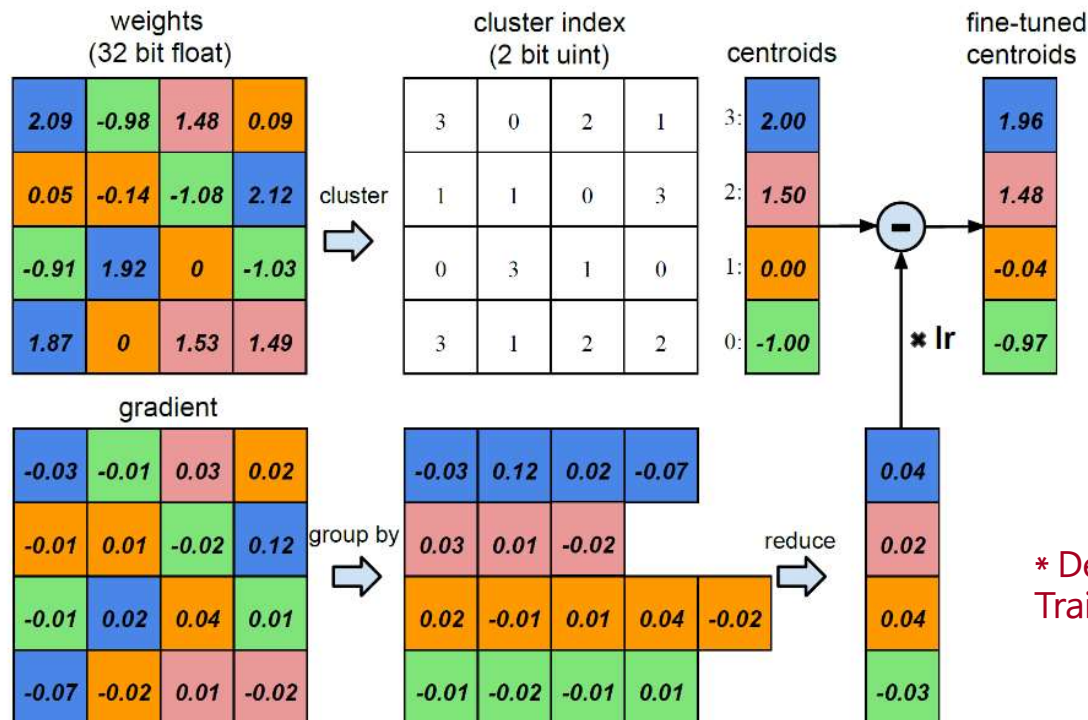


- Song Han, Compressing and regularizing deep neural networks, 2016
- Pavlo Molchanov et. al., "Pruning Convolutional Neural Networks for Resource Efficient Inference," ICLR 2017

Network	Top-1 Error	Top-5 Error	Parameters	Compression Rate
LeNet-300-100 Ref	1.64%	-	267K	
LeNet-300-100 Pruned	1.59%	-	22K	12×
LeNet-5 Ref	0.80%	-	431K	
LeNet-5 Pruned	0.77%	-	36K	12×
AlexNet Ref	42.78%	19.73%	61M	
AlexNet Pruned	42.77%	19.67%	6.7M	9×
VGG-16 Ref	31.50%	11.32%	138M	
VGG-16 Pruned	31.34%	10.88%	10.3M	13×

II. 모델 경량화 : Weight Sharing (1)

- Quantization and Weight Sharing
 - Cluster weights and use **centroid** of clustered weights in indexed array
 - Update centroid weights with **summation** of corresponding gradients
- Huffman Coding
 - Quantized weights are biased
 - Achieve additional compression via encoding weights



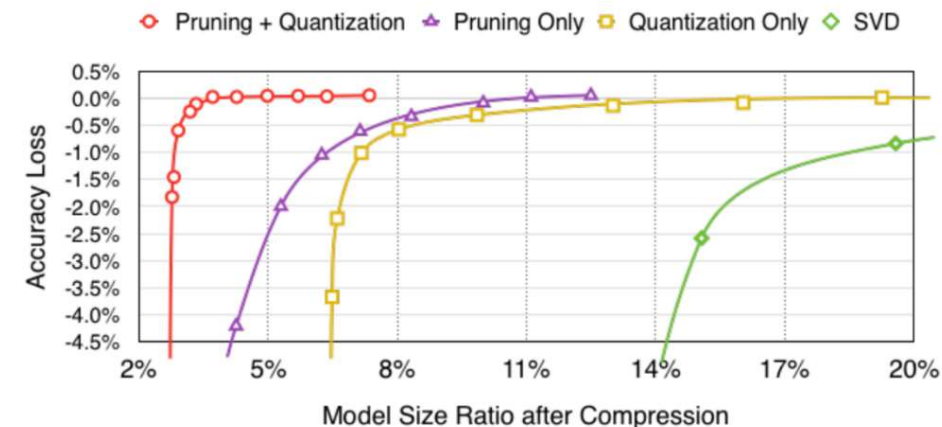
Biased weights

* Deep Compression: Compressing Deep Neural Networks with Pruning, Trained Quantization and Huffman Coding [Han et al., ICLR 2016]

II. 모델 경량화: Weight Sharing (2)

- Experiments
 - About 40x without critical accuracy loss
 - MNIST data with LeNet-300-100, LeNet-100
 - ImageNet data with AlexNet, VGG-16
- Does pruning and quantization make synergy?
 - Result becomes much better when using two methods together
 - Model can be compressed up to 3% of original size

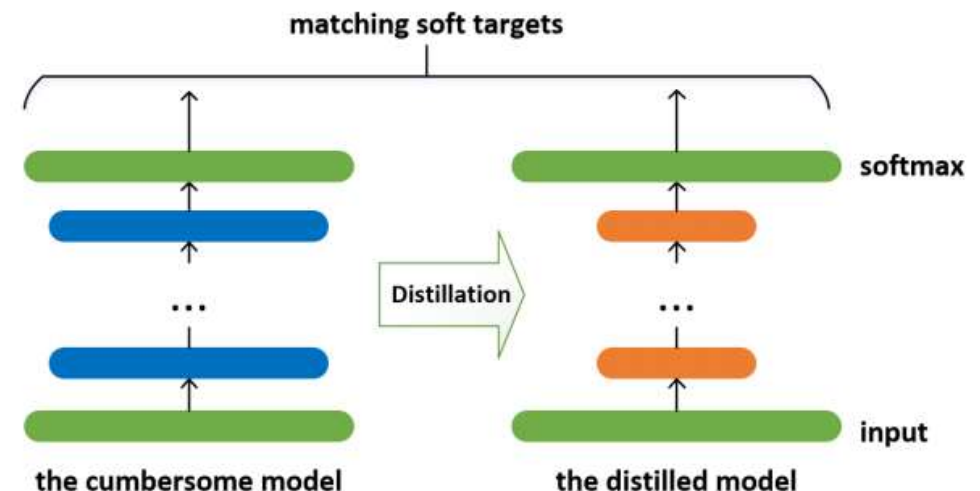

Network	Top-1 Error	Top-5 Error	Parameters	Compress Rate
LeNet-300-100 Ref	1.64%	-	1070 KB	
LeNet-300-100 Compressed	1.58%	-	27 KB	40×
LeNet-5 Ref	0.80%	-	1720 KB	
LeNet-5 Compressed	0.74%	-	44 KB	39×
AlexNet Ref	42.78%	19.73%	240 MB	
AlexNet Compressed	42.78%	19.70%	6.9 MB	35×
VGG-16 Ref	31.50%	11.32%	552 MB	
VGG-16 Compressed	31.17%	10.91%	11.3 MB	49×



[기존 모델의 가공(클러스터링) 통한 모델 경량화]

II. 모델 경량화: Distillation (1)

- Transfer the generalization ability of the cumbersome model to a small model
 - Use the class probabilities produced by the cumbersome model as “**soft targets**” for training the small model
 - Use the same training set or separate “**transfer set**” for the transfer stage
- **The Algorithm**
 1. Feed teacher with data (T_1)
 2. Obtain soft targets from teacher (T_1)
 3. Train student on soft targets (T_1) with cross-entropy loss
 4. Use student with $T < T_1$



II. 모델 경량화: Distillation (2)

- **Model Compression via Distillation and Quantization** [Polino et al., ICLR'18]
 - Given trained deep neural network (DNN), called 'teachers', make a compressed 'student' model, with similar accuracy using quantization and distillation
 - Teacher = Original Deep Model
 - Student = Quantized Model

II. 모델 경량화: Distillation (3)

- Quantization Process
 - Map each data point to the nearest quantization point.
 - Given quantization function $\hat{Q}()$,

$$Q(v) = \alpha \hat{Q} \left(\frac{v - \beta}{\alpha} \right) + \beta.$$

where $\alpha = \max_i v_i - \min_i v_i$, and $\beta = \min_i v_i$

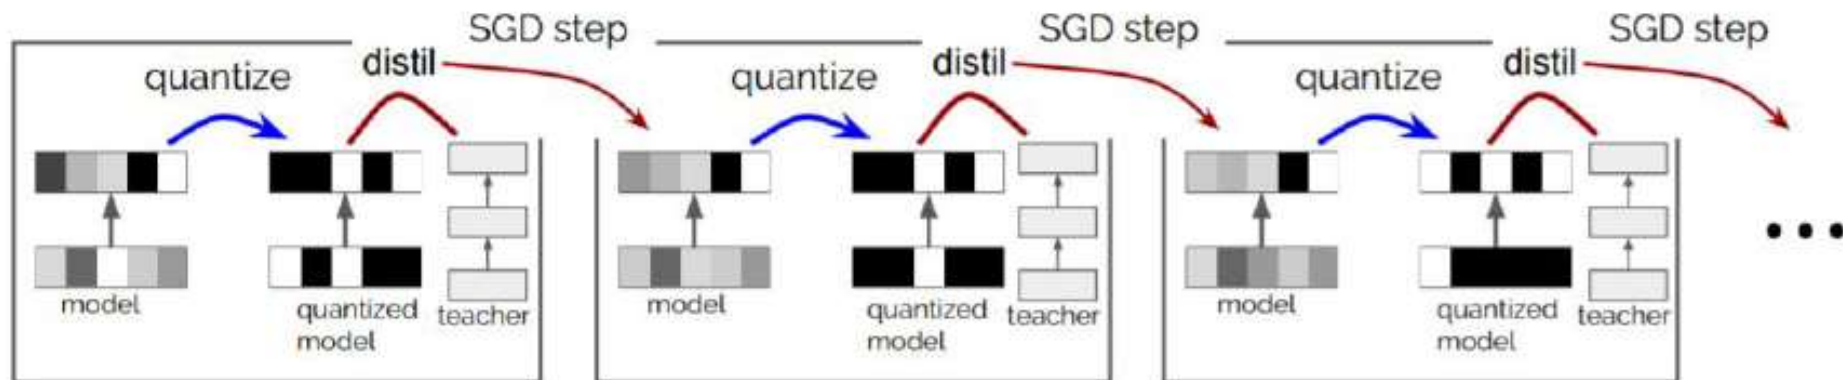
- The range of $\hat{Q}()$ function is $[0, 1]$
 - Example : if we quantize v in $[10, 50]$ into 5 quantization points, each v is mapped to the nearest point between $[10, 20, 30, 40, 50]$
 - Uniform vs non-uniform quantization
 - Quantization points can be learned from data

II. 모델 경량화: Distillation (4)

- Quantized Distillation

Algorithm 1 Quantized Distillation

```
1: procedure QUANTIZED DISTILLATION
2:   Let  $w$  be the network weights
3:   loop
4:      $w^q \leftarrow \text{quant\_function}(w, s)$ 
5:     Run forward pass and compute distillation loss  $l(w^q)$ 
6:     Run backward pass and compute  $\frac{\partial l(w^q)}{\partial w^q}$ 
7:     Update original weights using SGD in full precision  $w = w - \nu \cdot \frac{\partial l(w^q)}{\partial w^q}$ 
8:   Finally quantize the weights before returning:  $w^q \leftarrow \text{quant\_function}(w, s)$ 
9: return  $w^q$ 
```



II. 모델 경량화: Distillation (5)

CIFAR10 accuracy

		2 bits	4 bits	8 bits
Student model 1 1M param - 4 MB 84.5% - 88.8% X5.25	PM Quant.(No bucket)	9.30 %	67.99 %	88.91 %
	PM Quant. (with bucket)	10.53 %	87.18 %	88.80 %
	Quantized Distill.	82.4 %	88.00 %	88.82 %
	Differentiable Quant.	80.43%	88.31 %	—
Student model 2 0.3M param - 1.27 MB 80.3% - 84.3% X16.5	PM Quant. (No bucket)	10.15 %	68.05 %	84.38 %
	PM Quant. (with bucket)	11.89 %	81.96 %	84.38 %
	Quantized Distill.	74.22 %	83.92 %	84.22 %
	Differentiable Quant.	72.79 %	83.49 %	—
Student model 3 0.1M param - 0.45 MB 71.6% - 78.2% X46.66	PM Quant. (No bucket)	10.15 %	61.30 %	78.04 %
	PM Quant. (with bucket)	10.38 %	72.44 %	78.10 %
	Quantized Distill.	67.02 %	77.75 %	77.92 %
	Differentiable Quant.	57.84 %	77.36 %	—

Teacher model : 5.3M param., 21MB, accuracy 89.71%

OpenNMT dataset Bleu score and perplexity (ppl)

		2 bits	4 bits
Student model 1 81.6M param - 326 MB 14.97 - 16.13 BLEU	PM Quant.(No bucket)	0.00 - $2 \cdot 10^{17}$ ppl	0.24 - $2 \cdot 10^6$ ppl
	PM Quant. (with bucket)	4.12 - 125.1 ppl	16.29 - 26.2 ppl
	Quantized Distill.	0.00 - 6645 ppl	15.73 - 25.43 ppl
	Differentiable Quant.	0.7 - 249 ppl	15.01 - 28.8 ppl
Student model 2 64.8M param - 249 MB 14.22 - 15.48 BLEU	PM Quant. (No bucket)	0.00 - $5 \cdot 10^8$ ppl	6.65 - 71.78 ppl
	PM Quant. (with bucket)	1.72 - 286.98 ppl	15.19 - 28.95 ppl
	Quantized Distill.	0.00 - 4035 ppl	15.26 - 29.1 ppl
	Differentiable Quant.	0.28 - 306 ppl	13.86 - 31.33 ppl
Student model 3 57.2M param - 228 MB 12.45 - 13.8 BLEU	PM Quant. (No bucket)	0.00 - $3 \cdot 10^8$ ppl	5.47 - 106.5 ppl
	PM Quant. (with bucket)	0.24 - 1984 ppl	12.64 - 36.56 ppl
	Quantized Distill.	0.14 - 731 ppl	12 - 37 ppl
	Differentiable Quant.	0.26 - 306 ppl	12.06 - 38.44 ppl

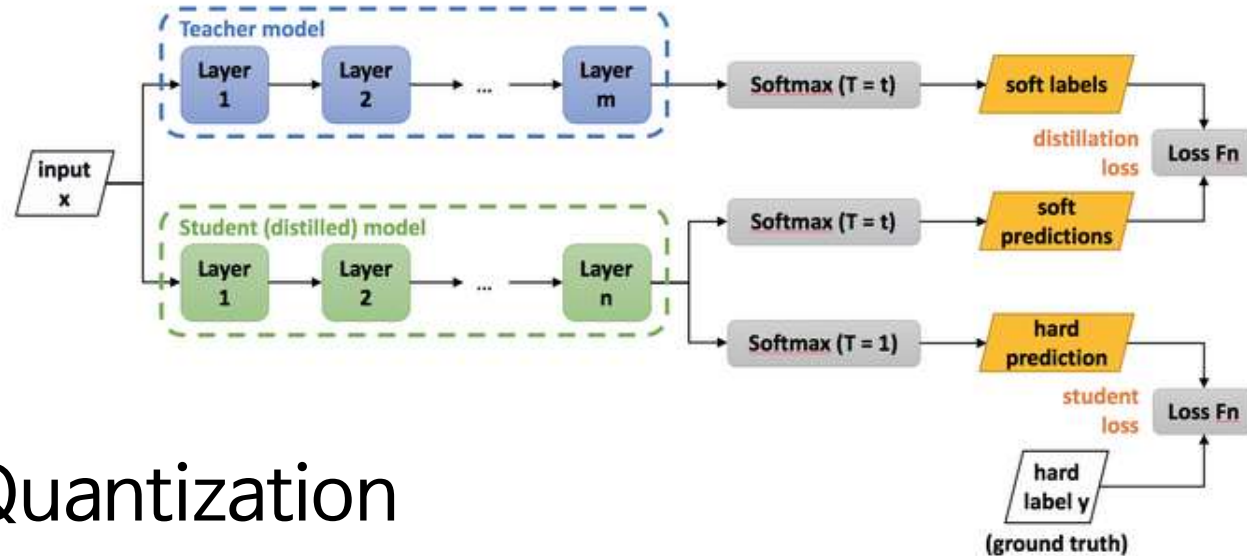
Teacher model : 84.8M
param., 340MB, 26.1 ppl
15.88 BLUE

[기존 모델에의 quantization과 Distillation을 통한 모델 경량화]

II. 모델 경량화: Distillation (6)

- Knowledge Distillation

[Kill-The-Bits'20]



- Product Quantization

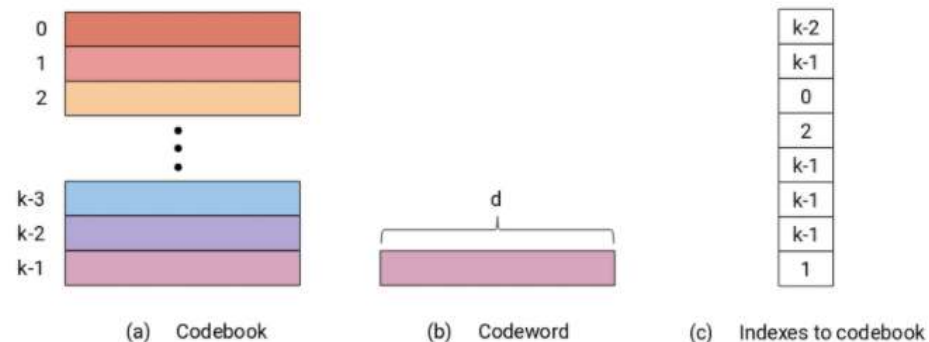
- Was utilized to quantize the weights of convolutional and fully-connected layer

$$\|y - \hat{y}\|_2^2 = \|x(W - q(W))\|_2^2$$

Objective function of quantization
for minimizing reconstruction error

VS

$$\|W - \hat{W}\|_2^2 = \sum \|w_j - q(w_j)\|_2^2$$



II. 모델 경량화: Distillation (7)

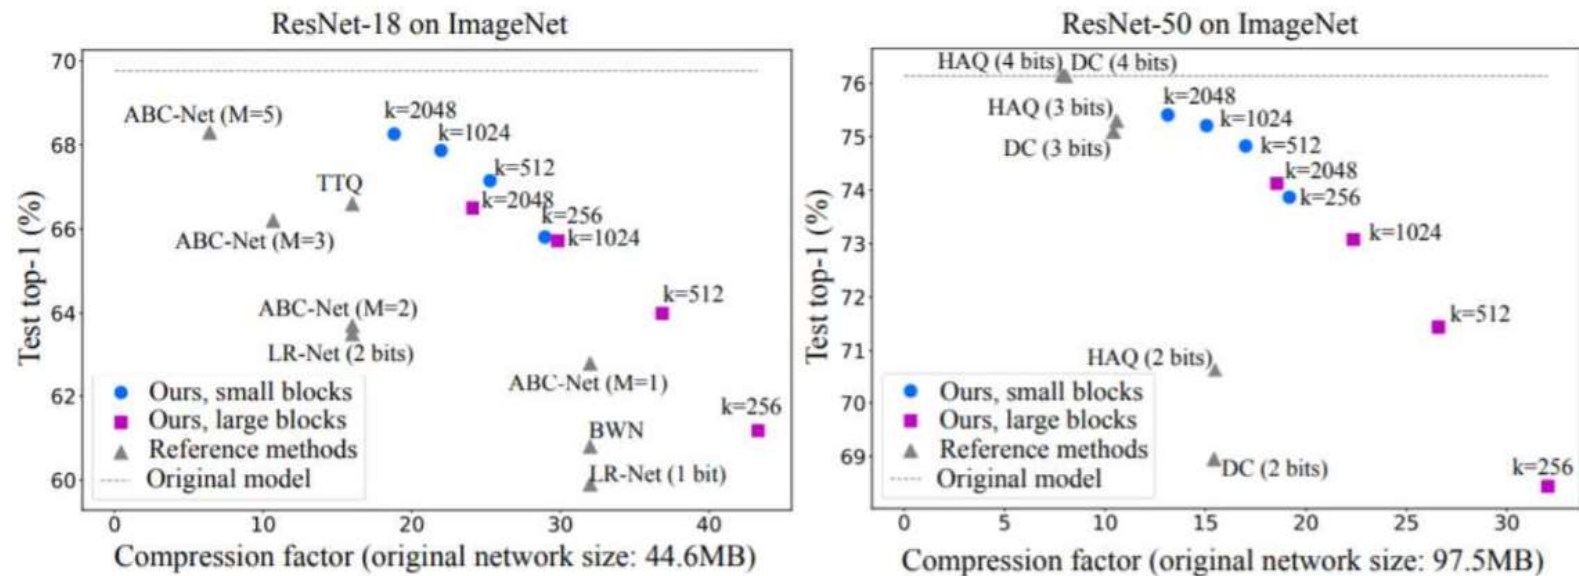


Figure 3: Compression results for ResNet-18 and ResNet-50 architectures. We explore two compression regimes as defined in Section 4.1: small block sizes (block sizes of $d=4$ and 9) and large block sizes (block sizes $d=8$ and 18). The results of our method for $k=256$ centroids are of practical interest as they correspond to a byte-compatible compression scheme.

Table 1: Results for vanilla ResNet-18 and ResNet-50 architectures for $k=256$ centroids.

Model (original top-1)	Compression	Size ratio	Model size	Top-1 (%)
ResNet-18 (69.76%)	Small blocks	29x	1.54 MB	65.81 ± 0.04
	Large blocks	43x	1.03 MB	61.10 ± 0.03
ResNet-50 (76.15%)	Small blocks	19x	5.09 MB	73.79 ± 0.05
	Large blocks	31x	3.19 MB	68.21 ± 0.04

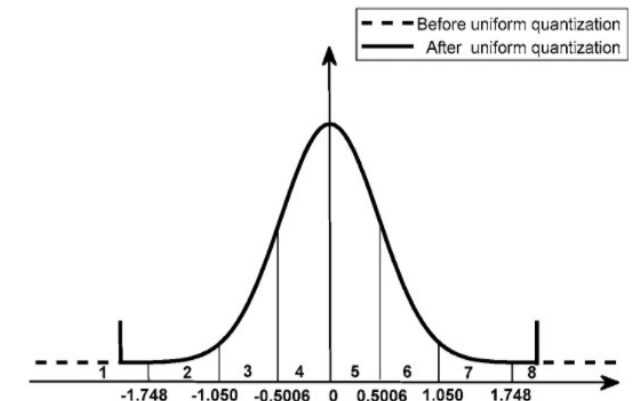


III. Quantization: Overview

- Main idea
 - Quantize FP32 typed weights, activations, or gradients to values with lower bit-widths
- Challenges
 - **Quantization errors** severely affect the precisions of DL models
 - Quantization makes **hard to converge** values of parameters
 - Specifically, values nearby boundaries of each quantization interval may be changed frequently
 - Quantization functions are **non-differentiable**
 - Use of alternative differentiable functions (e.g., STE)
 - Gradient mismatch problem

1.12	3.42	-1.5	-12
32	-1	-5	15
24	0.55	-54	0.24
-0.1	0.1	-0.2	2

→

1	1	-1	-1
1	-1	-1	1
1	1	-1	1
-1	1	-1	1



III. Quantization : Classification

- Methods

- **Quantization during training**

- BinaryConnect, Binary-weight, XNOR-Net, DoReFA-Net, LQ-Net, ...

- **Quantization after training**

* For more details, please refer to our survey paper,

E. Kim et al., *Communications of the KIISE* 38(8):18-29, Aug., 2020

- Codebook

Approaches	Types	codebooks	Representative work
Fixed codebook	<ul style="list-style-type: none">• Binarization• Scaled Binarization• Ternarization• Scaled Ternarization• Powers of two• K-bits	<ul style="list-style-type: none">• $\{-1, 1\}$• $\{-a, b\}$• $\{-1, 0, 1\}$• $\{-a, 0, b\}$• $\{0, \pm 1, \pm 2^{-1}, \dots, \pm 2^{-L}\}$• $\{\pm v1, \{\pm v1 \pm v2\}, \dots\}$	<ul style="list-style-type: none">• BinaryConnect, Binary-weight• XNOR-Net• Ternary Net• TTQ• [Tang and Kwan'94]• DoReFa-Net, LQ-Net
Adaptive codebook	<ul style="list-style-type: none">• Soft Quantization• Hard Quantization	<ul style="list-style-type: none">• Learned from data• Learned from data	<ul style="list-style-type: none">• Variational Network Quantization• Vector quantization, [Choi'2016]

- Targets for quantization

Components	Benefits	Challenges
Weights	<ul style="list-style-type: none">• Smaller model size• Faster forward training & inference• Less energy	<ul style="list-style-type: none">• Hard to converge with quantized weights• Require approximate gradients• Accuracy degradation
Activations	<ul style="list-style-type: none">• Smaller memory foot print during training• Allows replacement of dot-products by bitwise operations• Less energy	<ul style="list-style-type: none">• Gradient mismatch problem
Gradients	<ul style="list-style-type: none">• Communication & memory savings	<ul style="list-style-type: none">• Convergence requirement

III. Quantization : Binarization (1)

- Main idea

- Train DNNs with binary weights, while retraining precision of the stored weights in which gradients are accumulated to regularize all the parameters

- Binarization

- Deterministic

- $w_b = \begin{cases} +1 & \text{if } w \geq 0, \\ -1 & \text{otherwise.} \end{cases}$


- Stochastic

- $w_b = \begin{cases} +1 & \text{with probability } p = \sigma(w), \\ -1 & \text{with probability } 1 - p. \end{cases}$

where $\sigma(x) = \text{clip}(\frac{x+1}{2}, 0, 1) = \max(0, \min(1, \frac{x+1}{2}))$ is a hard sigmoid function which is used to limit p into $[0, 1]$.

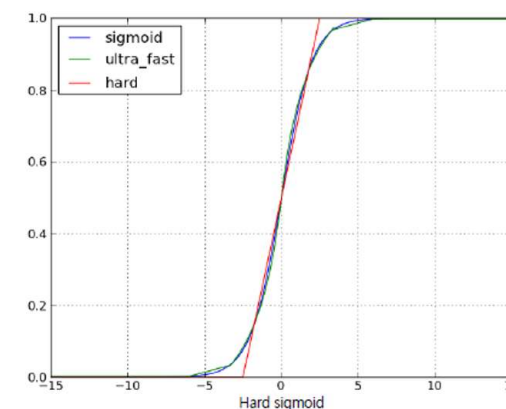
- Propagations & Updates

- w is binarized during forward and backward propagation
- Full-precision w is used during parameter update



1.12	3.42	-1.5	-12	1	1	-1	-1
32	-1	-5	15	1	-1	-1	1
24	0.55	-54	0.24	1	1	-1	1
-0.1	0.1	-0.2	2	-1	1	-1	1

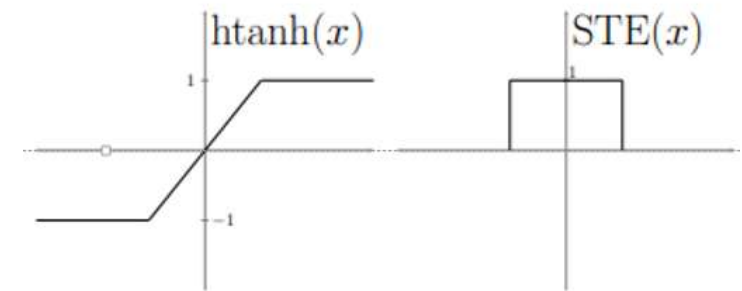
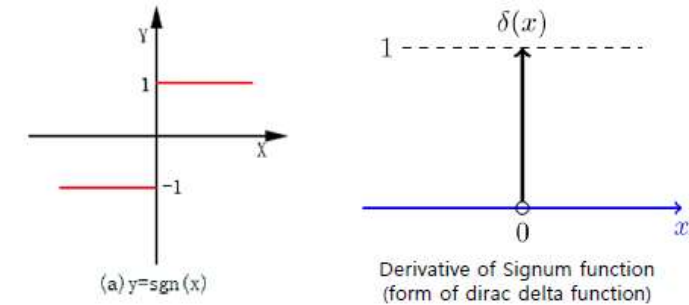
Binarization example



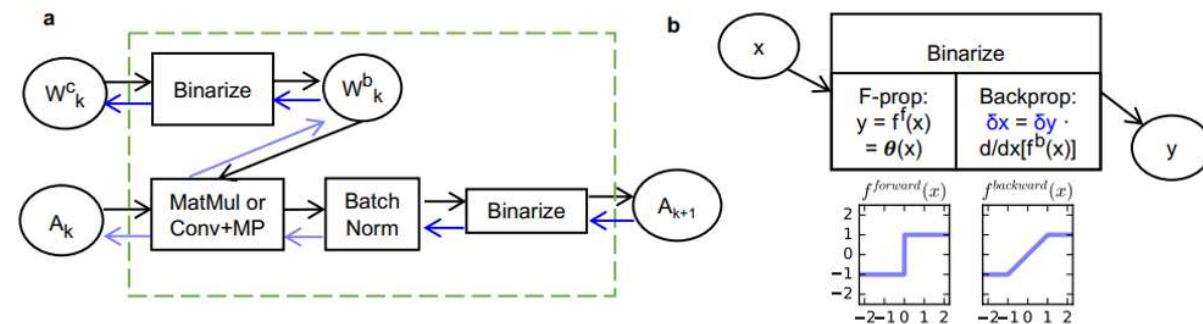
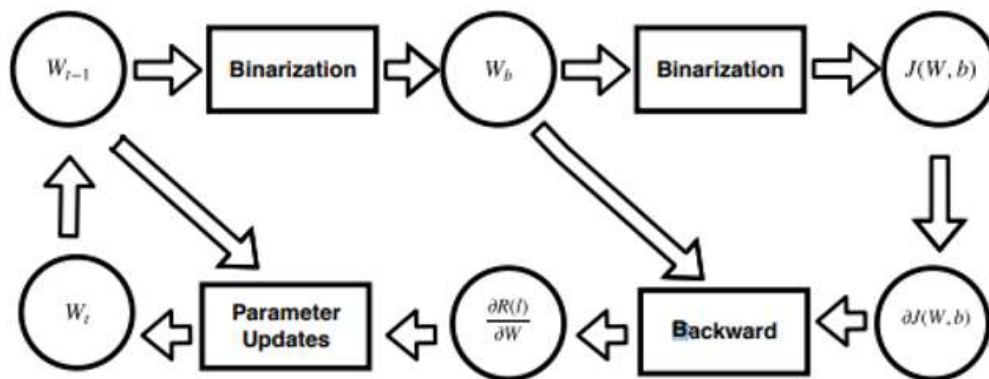
* Binarized Neural Networks: Training deep neural networks with weights and activations constrained to +1 or -1, [Courbariaux et al., 2016]

III. Quantization : Binarization (2)

- Forward & Backward propagation
 - The derivative of the sign function is zero almost everywhere \rightarrow Gradient Vanishing
 - How to back-propagate gradient through ?
 - “straight-through estimator(STE)”, previously introduced by Hinton (2012)




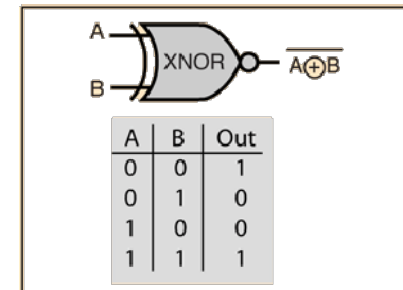
Procedure



Binary convolution transformer(dashed green line)

III. Quantization : XNOR-NET (1)

- Goal: Simple but efficient approximations to CNN by binarizing the weights and intermediate representations in CNN
 - Convolutions can be estimated by only addition and subtraction (X2 speedup)

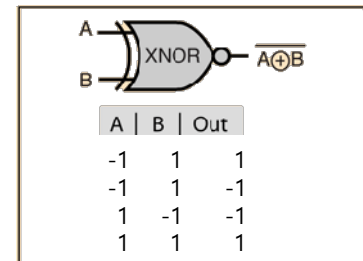


	Network Variations	Operations used in Convolution	Memory Saving (Inference)	Computation Saving (Inference)	Accuracy on ImageNet (AlexNet)
Standard Convolution	Real-Value Inputs $\begin{bmatrix} 0.11 & -0.21 & \dots & -0.34 \\ -0.25 & 0.61 & \dots & 0.52 \end{bmatrix}$ Real-Value Weights $\begin{bmatrix} 0.12 & -1.2 & 0.41 \\ -0.2 & 0.5 & 0.68 \end{bmatrix}$	$+, -, \times$	1x	1x	%56.7
Binary Weight	Real-Value Inputs $\begin{bmatrix} 0.11 & -0.21 & \dots & -0.34 \\ -0.25 & 0.61 & \dots & 0.52 \end{bmatrix}$ Binary Weights $\begin{bmatrix} 1 & -1 & 1 \\ -1 & 1 & 1 \end{bmatrix}$	$+, -$	$\sim 32x$	$\sim 2x$	%56.8
BinaryWeight Binary Input (XNOR-Net)	Binary Inputs $\begin{bmatrix} 1 & -1 & \dots & -1 \\ -1 & 1 & \dots & 1 \end{bmatrix}$ Binary Weights $\begin{bmatrix} 1 & -1 & 1 \\ -1 & 1 & 1 \end{bmatrix}$	XNOR, bitcount	$\sim 32x$	$\sim 58x$	%44.2


III. Quantization : XNOR-NET (2)

$$\mathbf{C}^* = \text{sign}(\mathbf{Y}) = \text{sign}(\mathbf{X}^T) \text{sign}(\mathbf{W}) = \mathbf{H}^{*\top} \mathbf{B}^*$$


- Binarize both weights and inputs (previous results)
- Convolution as **Binary dot product** implemented by XNOR-BitCounting operations



(1) Binarizing Weight

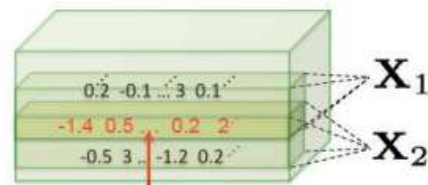


$$\frac{1}{n} \|\mathbf{W}\|_{\ell_1} = \alpha$$



(2) Binarizing Input

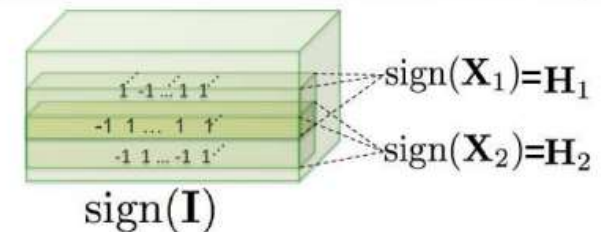
Inefficient



Redundant computations in overlapping areas


$$\frac{1}{n} \|\mathbf{X}_1\|_{\ell_1} = \beta_1$$

$$\frac{1}{n} \|\mathbf{X}_2\|_{\ell_1} = \beta_2$$



(3) Binarizing Input

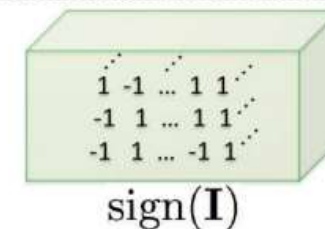
Efficient





$$\mathbf{A} * \mathbf{k} = \mathbf{K}$$

$$\beta_1$$

$$\beta_2$$



(4) Convolution with XNOR-Bitcount



$$\mathbf{I} * \mathbf{W} \approx \left[\mathbf{sign}(\mathbf{I}) * \mathbf{sign}(\mathbf{W}) \right] \odot \mathbf{K} \odot \alpha$$

[기존 연산의 bitwise 연산자로 치환 가능성, Backward Prop.를 위한 Real weight 유지 필요]

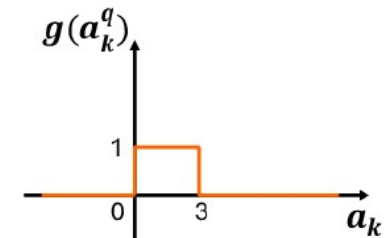
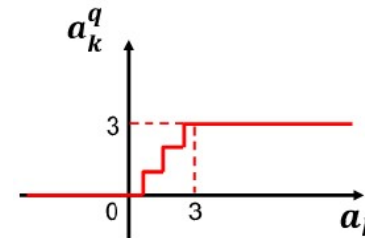
III. Quantization : DoReFa-Net

- Goal: Forward/backward passes operate on low bitwidth **weights, activations** and **gradients**
 - Gradients are stochastically quantized before being propagated to convolutions
- Multi-bit quantization

An STE we will use extensively in this work is **quantize_k** that quantizes a real number input $r_i \in [0, 1]$ to a k -bit number output $r_o \in [0, 1]$. This STE is defined as below:

$$\textbf{Forward: } r_o = \frac{1}{2^k - 1} \text{round}((2^k - 1)r_i) \quad (5)$$

$$\textbf{Backward: } \frac{\partial c}{\partial r_i} = \frac{\partial c}{\partial r_o}. \quad (6)$$



* DoReFa-Net: Training Low Bitwidth Convolutional Neural Networks with Low Bitwidth Gradients, S. Zhou et al., arXiv:1606.06160

[LQ-Nets'18]

- Existing methods often use **simple, hand-crafted** quantizers
 - e.g., uniform or logarithmic quantization
- Learnable quantizer
 - Optimal quantizer should yield minimal **quantization error** for the input data distribution

$$Q^*(x) = \arg \min_Q \int p(x)(Q(x) - x)^2 dx,$$

- Distributions can be complex and different at different layers

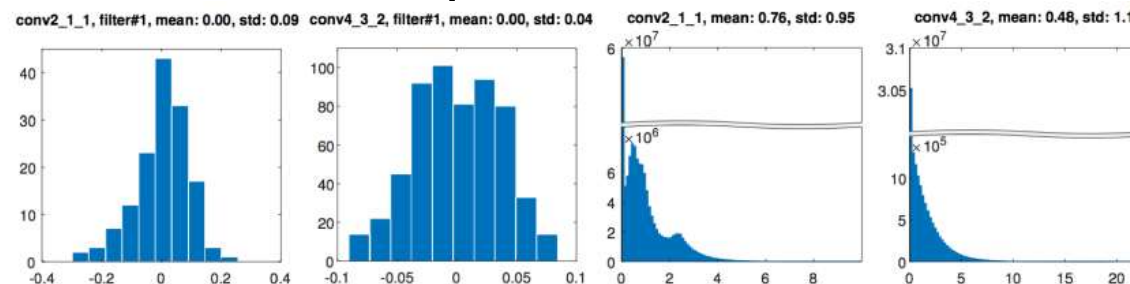



Fig. 1: Distributions of weights (left two columns) and activations (right two columns) at different layers of the ResNet-20 network trained on CIFAR-10. All the test-set images are used to get the activation statistics.

- Bitwise operator compatibility
 - Training quantizers to optimize the quantization level $\{q_i\}$ hampers bitwise operation compatibility
 - Solution is to separate mappings between floating-point basis and bits and learn the basis
 - Confine quantizations into subspaces compatible with bitwise operations



- An integer q represented by a K -bit binary encoding is the inner product btw. A basis vector and the binary coding vector \mathbf{b}

$$q = \left\langle \begin{bmatrix} 1 \\ 2 \\ \vdots \\ 2^{K-1} \end{bmatrix}, \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_K \end{bmatrix} \right\rangle, \quad \mathbf{b} = [b_1, b_2, \dots, b_K]^T \text{ where } b_i \in \{0, 1\}$$


└ Basis vector with K scalars

Table 4: Impact of bit-width on our LQ-Nets

ResNet-20 (CIFAR-10)		VGG-Small (CIFAR-10)		ResNet-18 (ImageNet)	
Bit-width (W/A)	Acc. (%)	Bit-width (W/A)	Acc. (%)	Bit-width (W/A)	Acc. (%)
32/32	92.1	32/32	93.8	32/32	70.3
1/32	90.1	1/32	93.5	2/32	68.0
2/32	91.8	2/32	93.8	3/32	69.3
3/32	92.0	3/32	93.8	4/32	70.0
1/2	88.4	1/2	93.4	1/2	62.6
2/2	90.2	2/2	93.5	2/2	64.9
2/3	91.1	2/3	93.8	3/3	68.2
3/3	91.6	3/3	93.8	4/4	69.3

III. Quantization: SketchML (1)

- **Goal** : Build a compression method that can efficiently handle a sparse & nonuniform gradients consisting of key-value pairs in distributed ML
- Distributed ML
 - Data-parallelism
 - Data split over multiple machines
 - Model replicas train over different parts of data & communicate model information periodically
 - Averaging gradients from workers
 - Model-parallelism
 - Models split over multiple machines
 - A single training iteration spans multiple machines
 - **Reducing the size of gradients to be communicated** is a major challenge



Parameter server [OSDI'14]

III. Quantization: SketchML (2)

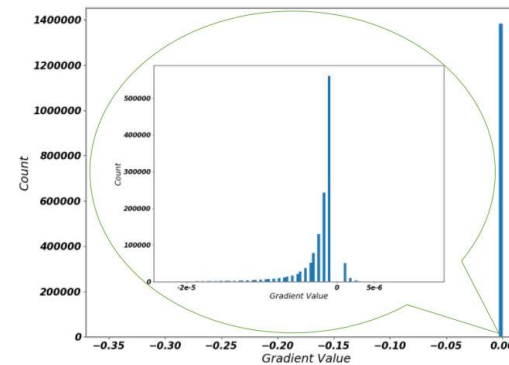
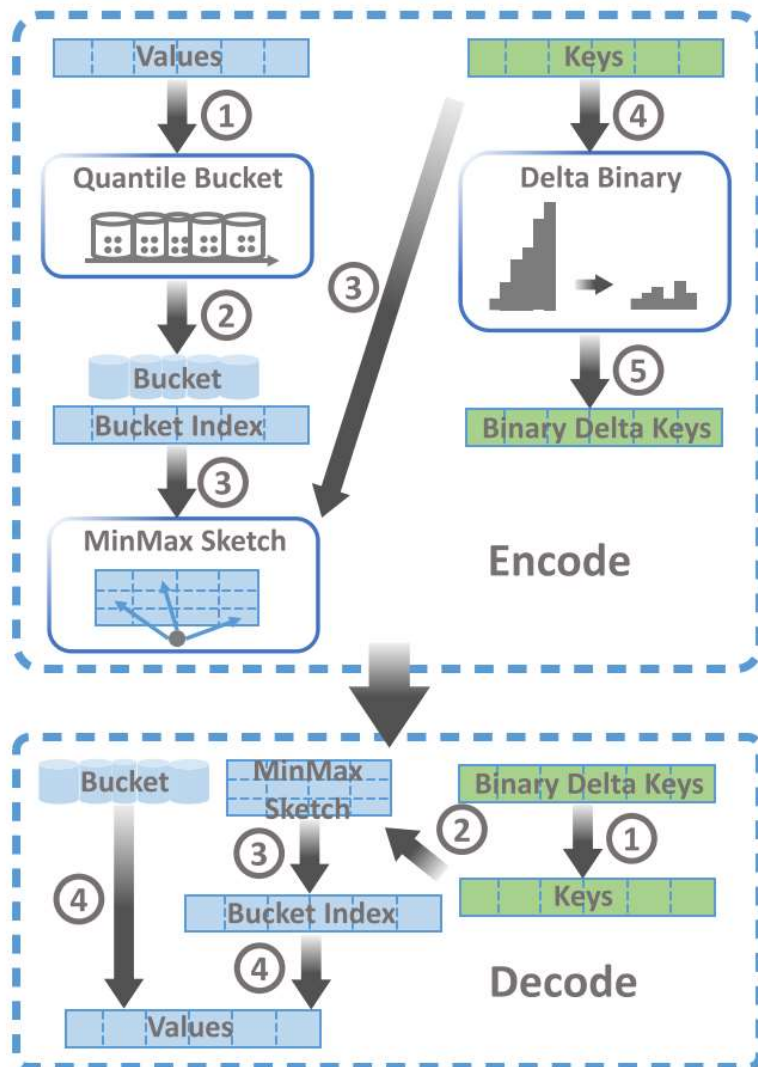
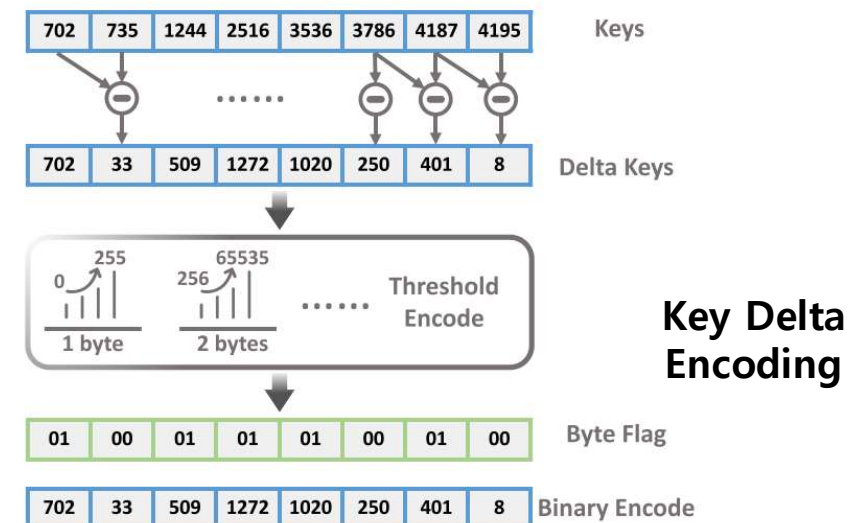
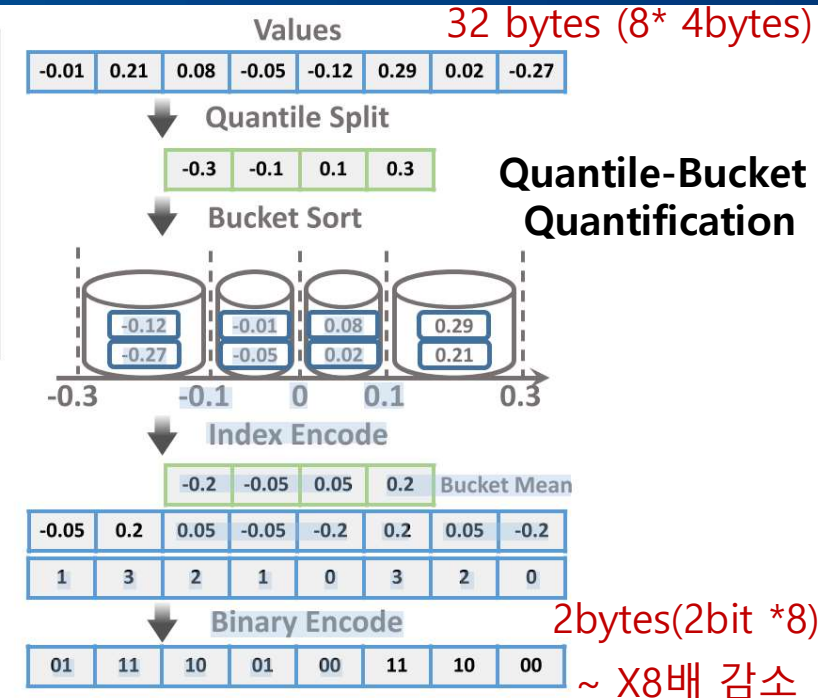
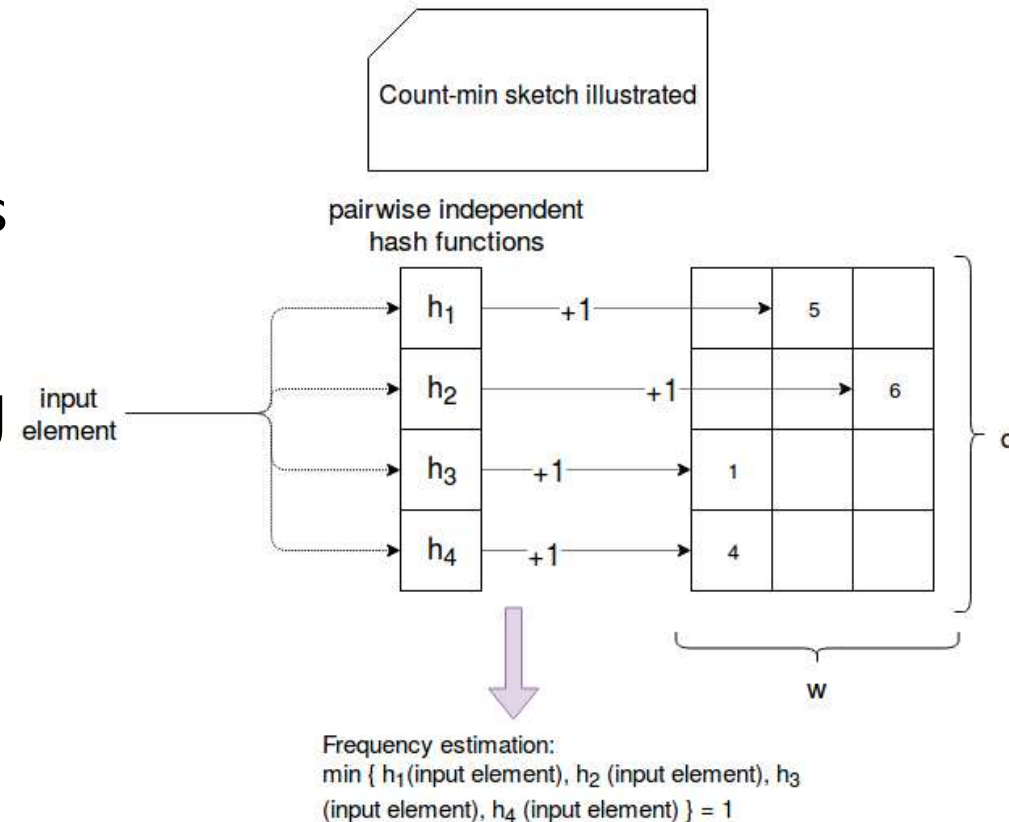


Figure 4: Nonuniform Gradient Values



Appendix: Sketch algorithm overview

- Idea:
 - Summarize your data with a **probabilistic** data structure
- **Heavy hitters** problem
 - focuses on the retrieval of all elements appearing at least $x\%$ times in a given data stream
- We can also classify the use of Count-min sketch in the following categories of queries:
 - **Point query**: retrieves the estimated number of occurrences for one particular event
 - **Inner product query**: computes the inner product of 2 vectors.; useful to estimate the join size in relational query processing
 - **Range query**: counts the sum of elements between 2 range values



IV. Compact Network Design (1)

- Problem #1 : Extensive cost
 - With more channels, we can learn more filters
 - However, # of parameters increases as we have more channels.
- Problem #2 : Dead channels
 - Some channels do not affects the outputs at all.
- Problem #3 : Low correlation btw. Channels

IV. Compact Network Design (1)


- Standard Convolution

- # of Parameters

- K : filter size, C : # of Input channels, M : # of output channels
 - One filter size : $K^2 C$
 - Total # of parameters : $K^2 CM$

- Computational cost

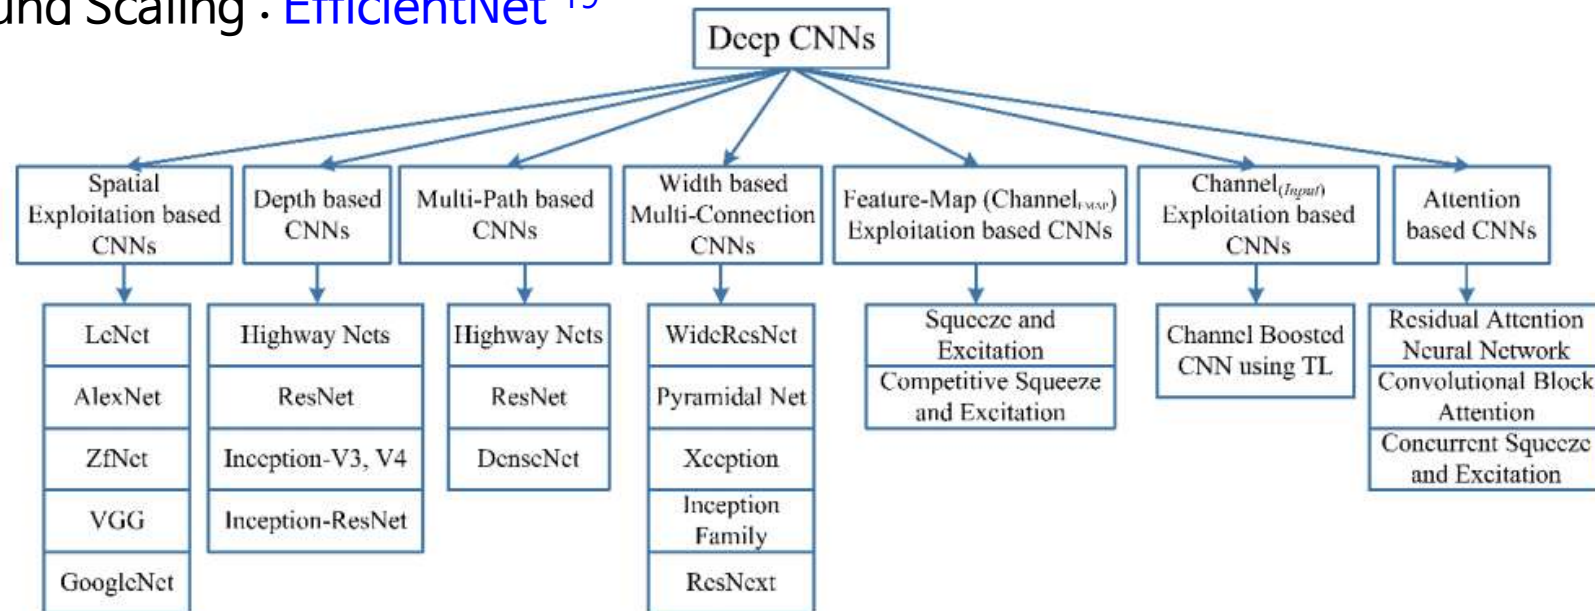
- Input size : $H \times W$, output size : $H \times W$
 - $K^2 CMHW$



- CNN 특징 : Full connectivity가 아닌 local connectivity로 계산량 감소
- 최적화 요소 : Activation, Pooling, Architecture, Regularization, Optimization 등

IV. Compact Network Design (2)

- Design optimized neural network architecture
 - Grouped Convolution : AlexNet^{'12}, ShuffleNet^{'17}
 - Residual connection, Bottleneck : Inception^{'15}, ResNet^{'16}
 - Depth-wise Conv., Point-wise conv.: Inception^{'15}, MobileNet^{'17}
 - Dilated Convolution: DeepLab^{'18}, ShiftNet^{'17~'19}
 - Shift Convolution: ShiftNet^{'17~'19}
 - Compound Scaling : EfficientNet^{'19}

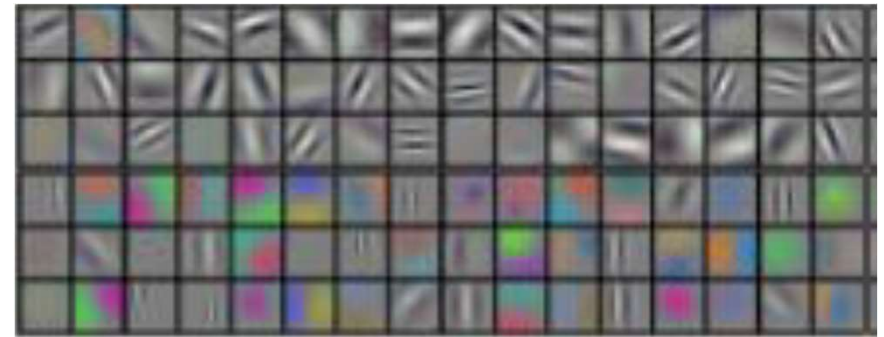


* Khan, Asifullah, et al. "A survey of the recent architectures of deep convolutional neural networks." *Artificial Intelligence Review* (2020): 1-62.

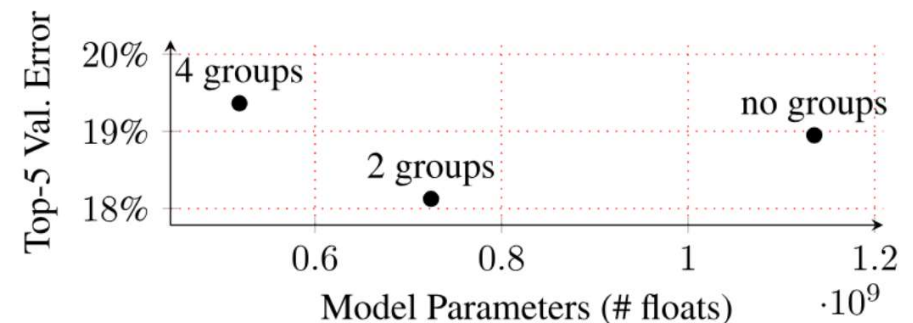
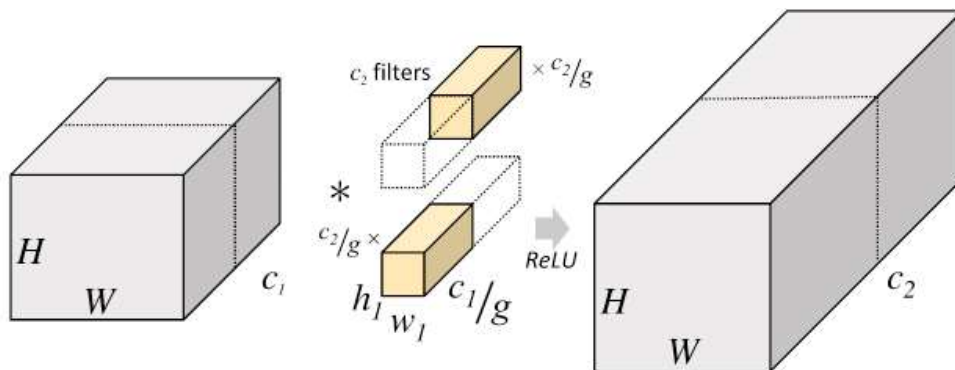
Grouped Convolution

- Group filters to separate channel information and learn (benefit to parallel processing)
- Sparse by learning highly correlated information for each filter group (fewer parameters)
- # of Parameters
 - $(K^2CM)/g$; g : # of groups
- Computational cost
 - $(K^2CMHW)/g$

[AlexNet, ShuffleNet]

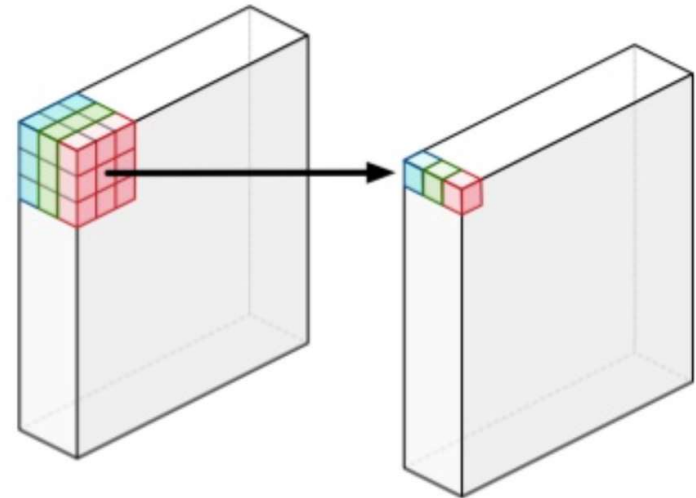


- (Top) Shading-originated kernel learning group
- (Bottom) Kernel learning group focused on colors & patterns




Depthwise Convolution

- Channel-independent convolution
 - Standard conv. Is impossible to spatial features for each channel
 - Each filter works for each channel
- # of parameters
 - K^2C
- Computational cost
 - K^2CMW



Point-wise convolution


- Effects of dimensional reduction by reducing # of channels
- Use of 1x1 convolutional filter
- # of parameters
 - CM since $K = 1$
- Computational cost
 - $CMHW$ since $K = 1$



DepthWise Separable PointWise Conv.

- Xception*, MobileNet**
 - CNN 계산량 및 매개변수 수를 줄임으로써 Neural Network 경량화
 - C 채널수, (X, Y) 입력 영상 크기, K 커널 필터 크기, M 필터 출력 채널 수

	Spatial Convolution	Depthwise Separable Conv.
operation 수	MCK^2	$C(K^2 + M)$
Parameter 수	$CHWK^2M$	$CHWK^2 + CHWM$ $= CHW(K^2 + M)$



* Francois Chollet, "Xception: Deep Learning with Depthwise Separable Convolutions," CVPR 2017

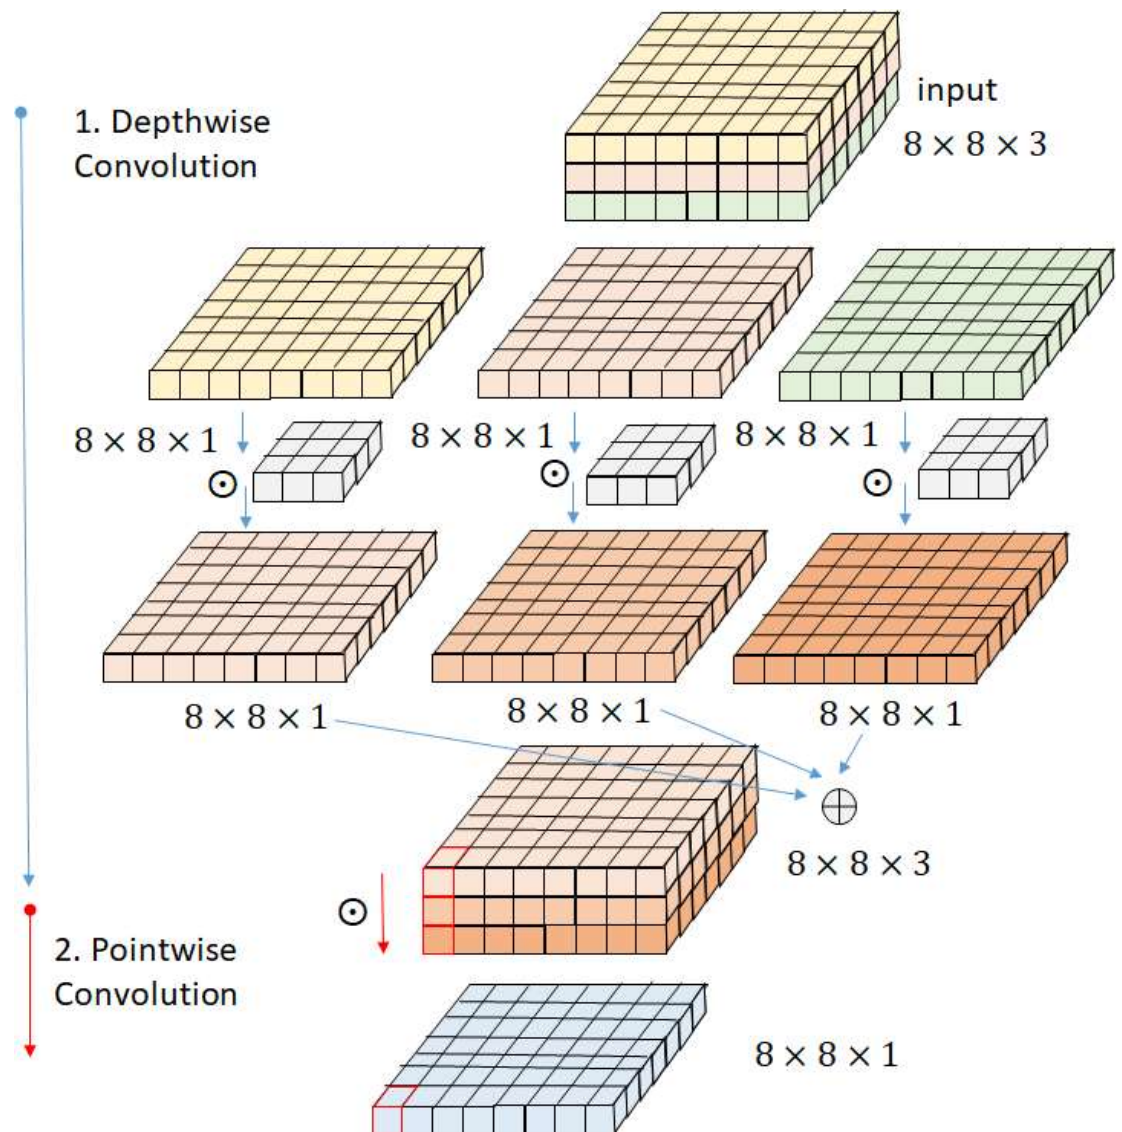


** Andrew G. Howard et al., "MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications," CVPR 2017

Reinterpretation of Convolutions

[MobileNet, Inception]

Depth-Wise Separable Point-wise Convolution


General Convolution



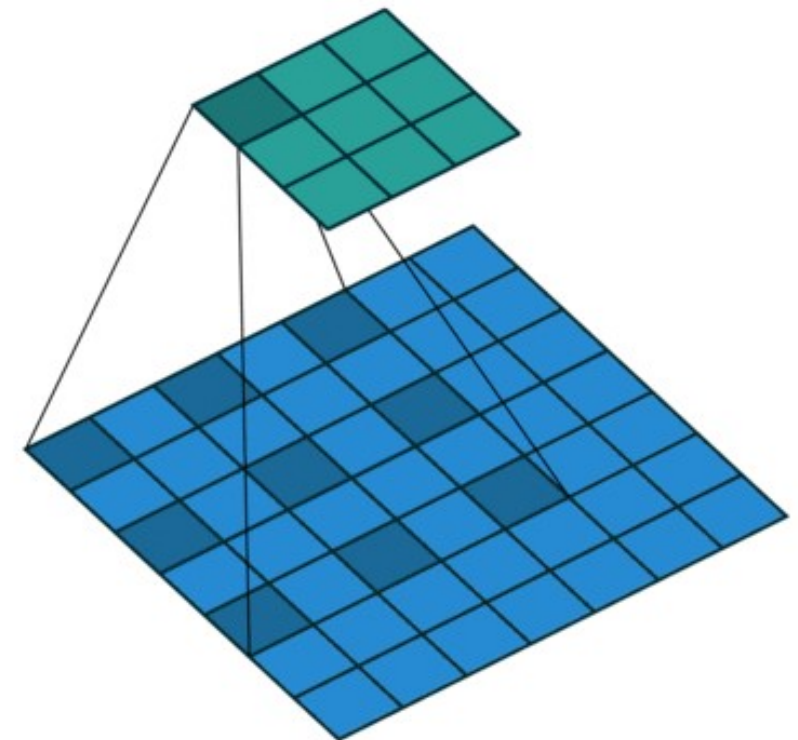
- Overview

[DeepLab, ShiftNet]

- Receptive field area is widened, but by maintaining the kernel size
- # of convolution operations is maintained while obtaining a pooling effect to improve accuracy



r : dilation rate, the spacing between kernels



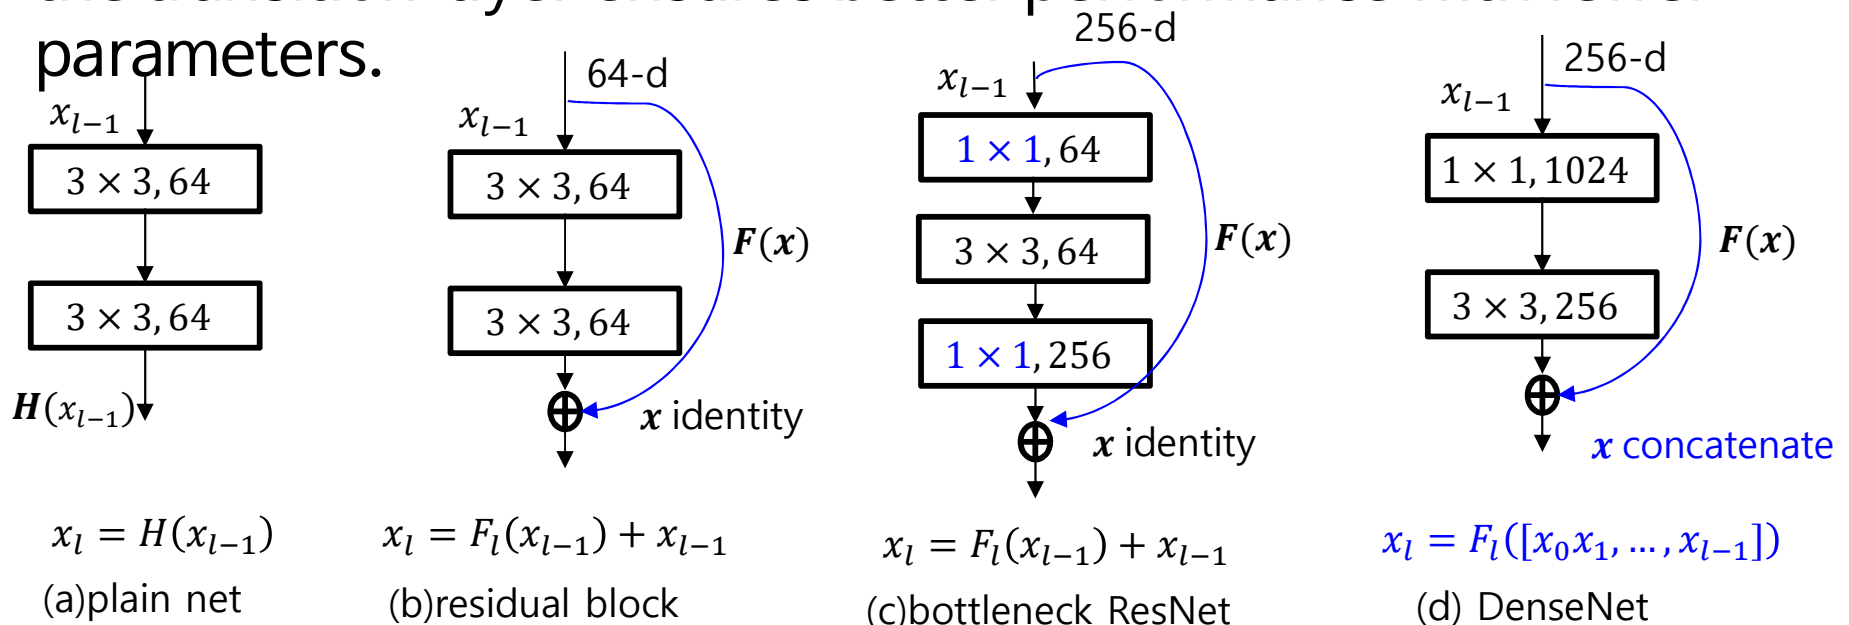
[ResNet, DenseNet]

• Residual Connection


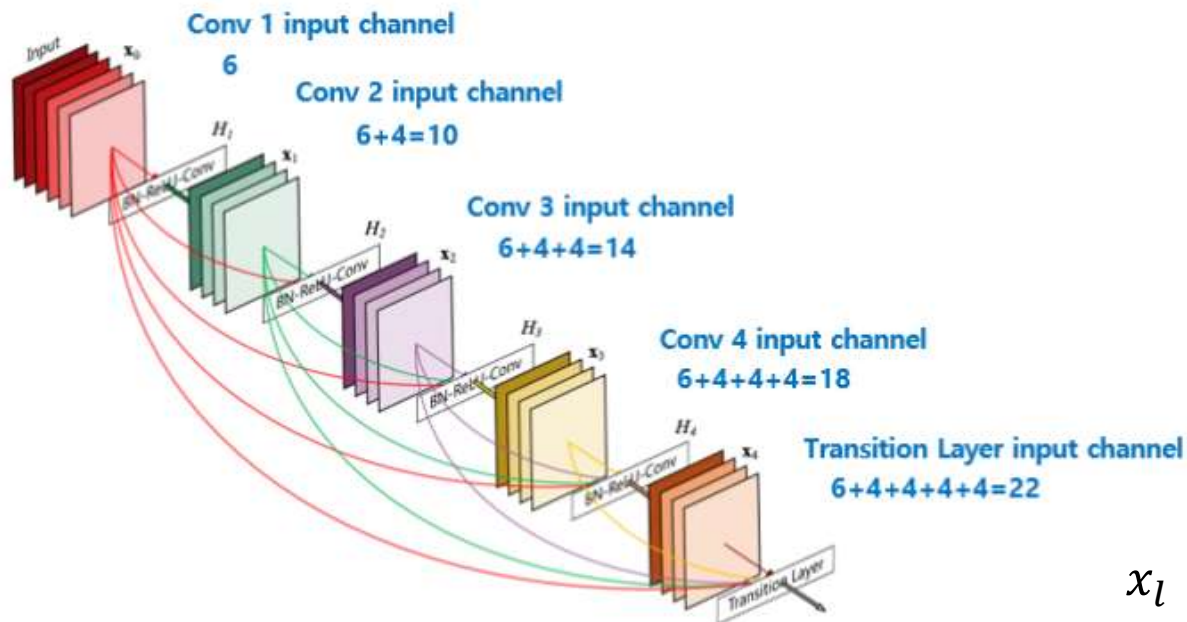
- With the same number of parameters, there is little additional computation other than addition, so deep networks are well optimized.

• Dense Net


- While collecting feature maps as concatenation, compression at the transition layer ensures better performance with fewer parameters.



- Dense Convolution (DenseNet*)
 - 기존의 feature map을 additive 형태가 아닌 concat 형태로 취합
 - feature-maps 압축으로 적은 parameter로 더 나은 성능 확인



bottleneck
(for ResNet)



bottleneck
(for DenseNet)

$$x_l = H_l(x_{l-1}) + x_{l-1} \quad x_l = H_l([x_0, x_1, \dots, x_{l-1}])$$

* G. Huang et al., "Densely Connected Convolutional Networks," CVPR 2017

- Shift (ShiftNet*, **)
 - Spatial convolution을 shift operation으로 대체하여 연산을 줄임




Figure 2: Illustration of (a) spatial convolutions, (b) depth-wise convolutions and (c) shift. In (c), the 3x3 grids denote a shift matrix with a kernel size of 3. The lighted cell denotes a 1 at that position and white cells denote 0s.





Figure 1: Illustration of a shift operation followed by a 1x1 convolution. The shift operation adjusts data spatially and the 1x1 convolution mixes information across channels.

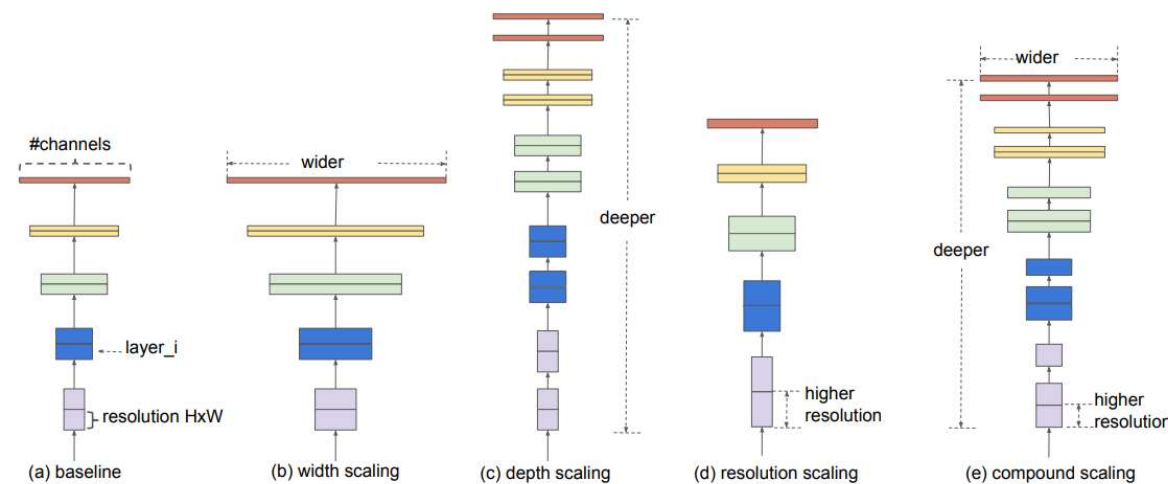


* B. Wu et al., "Shift: A zero FLOP, Zero Parameter Alternative to Spatial Convolutions," CVPR 2017


** W. Chen et al., "All You Need is a Few Shifts: Designing Efficient Convolutional Neural Networks for Image Classification," CVPR 2019

• Compound Scaling (EfficientNet*)

- CNN model의 3요소 (Depth α , Width β , Resolution γ)의 Efficient한 Scaling 방법에 대한 연구
- mNasNet**의 결과 모델을 base model로 찾음
- 최종적으로 $\alpha = 1.2$, $\beta = 1.1$, $\gamma = 1.15$ 를 고정한 채, ϕ 를 늘려 $B_1 \sim B_7$ 모델 탐색



Stage i	Operator \hat{F}_i	Resolution $\hat{H}_i \times \hat{W}_i$	#Channels \hat{C}_i	#Layers \hat{L}_i
1	Conv3x3	224×224	32	1
2	MBConv1, k3x3	112×112	16	1
3	MBConv6, k3x3	112×112	24	2
4	MBConv6, k5x5	56×56	40	2
5	MBConv6, k3x3	28×28	80	3
6	MBConv6, k5x5	28×28	112	3
7	MBConv6, k5x5	14×14	192	4
8	MBConv6, k3x3	7×7	320	1
9	Conv1x1 & Pooling & FC	7×7	1280	1



* Mingxing Tan, Quoc V. Le, "Efficient Net: Rethinking Model Scaling for Convolutional Neural Networks," ICML 2019

** Mingxing Tan et. al., "MnasNet: Platform-Aware Neural Architecture Search for Mobile," CVPR 2018

• Compound Scaling (EfficientNet)

- 타모델 대비 Efficient Net은 parameter수를 8~9배 줄이면서 정확도 향상
- Inference 속도 성능은 6배 정도 향상





Table 4. Inference Latency Comparison

Acc. @ Latency	
ResNet-152	77.8% @ 0.554s
EfficientNet-B1	78.8% @ 0.098s
Speedup	5.7x
Acc. @ Latency	
GPipe	84.3% @ 19.0s
EfficientNet-B7	84.4% @ 3.1s
Speedup	6.1x

Figure 1. Model Size vs. ImageNet Accuracy. All numbers are

Figure 5. FLOPS vs. ImageNet Accuracy – Similar to Figure 1

* Mingxing Tan, Quoc V. Le, "Efficient Net: Rethinking Model Scaling for Convolutional Neural Networks," ICML 2019